

HW3: Stereo Imaging

CS484 Introduction to Computer Vision

Homework 3 supplementary slides




Carefully look all materials
(webpage, code and supplementary slides)

Filter Demosaic



- Demosaic the raw image using the following three methods
 1. Down-sampling
 2. Bilinear interpolation
 3. Bicubic interpolation



Bayer pattern
(RGGB)



Raw image



Color image
(reference)



Demosaic image

(Optional) Bicubic Interpolation



- Bicubic interpolation (+10 pts)


$$f(x, y) = \sum_{i=0}^3 \sum_{j=0}^3 a_{ij} x^i y^j \quad x, y \in [0, 1] \times [0, 1]$$

- 16 unknown coefficients a_{ij}
- 16 known equations
 - $f(0,0), f(0,1), f(1,0), f(1,1)$
 - $f_x(0,0), f_x(0,1), f_x(1,0), f_x(1,1)$
 - $f_y(0,0), f_y(0,1), f_y(1,0), f_y(1,1)$
 - $f_{xy}(0,0), f_{xy}(0,1), f_{xy}(1,0), f_{xy}(1,1)$
- Difference approximation

$$f_x(x, y) = [f(x+1, y) - f(x-1, y)] / 2$$

$$f_y(x, y) = [f(x, y+1) - f(x, y-1)] / 2$$

$$f_{xy}(x, y) = [f(x+1, y+1) + f(x-1, y-1) - f(x+1, y-1) - f(x-1, y+1)] / 4$$




Epipolar geometry



World coordinate \mathbf{X} projects to image coordinate \mathbf{x} and \mathbf{x}'

What is the relation between \mathbf{x} and \mathbf{x}' ?




Epipolar geometry



The camera centers C and C' , a 3D point X , and its image x and x' lie in a common plane π .

The plane π is **epipolar plane**.




Epipolar geometry



World coordinate \mathbf{X} projects to image coordinate \mathbf{x} , but it can't distinguish with dots on the ray from \mathbf{C} to \mathbf{X} .


The projection of the ray from \mathbf{C} to \mathbf{X} on the image plane 2 is the line \mathbf{l}' .



Epipolar geometry




- The line l' is the **epipolar line**.
- The projection of X should be on the line l' .
- It is also the intersection of epipolar plane and image plane.



Epipolar geometry



- Intersection of the epipolar planes is **baseline**.
- C projects to e' , that every epipolar line cross. The point e' is **epipole**.
- The epipole is the intersection of the baseline and the image plane.
- Every epipolar line intersect on the epipole.



Fundamental matrix




- We want to know the relation between \mathbf{x} and \mathbf{l}' .
- The line \mathbf{l}' can be represented by

$$a'x' + b'y' + c' = 0$$

- \mathbf{l}' can be define as

$$\mathbf{l}' = \begin{bmatrix} a' \\ b' \\ c' \end{bmatrix} \quad \mathbf{l}'^T \mathbf{x}' = \mathbf{x}'^T \mathbf{l}' = 0$$

- The scale of \mathbf{l}' can be changed.
(i.e., $k\mathbf{l}'$ for $k \neq 0$ indicates the identical line)




Fundamental matrix



- Line l' passes through points x' and e' .
- l' is perpendicular to both x' and e' (as vectors in \mathbb{R}^3 , just numerically)
- Thus, l' can be written as $l' = e' \times x'$
- Cross product can be represented by multiplication with a skew-symmetric matrix


$$[e']_{\times} = \begin{bmatrix} 0 & -e'_3 & e'_2 \\ e'_3 & 0 & -e'_1 \\ -e'_2 & e'_1 & 0 \end{bmatrix} \quad e' \times x' = [e']_{\times} x'$$



Fundamental matrix



- \mathbf{x} can project to **any** plane π . The projected point is \mathbf{X} .
- The transformation from a 2D plane to another 2D plane is homography.
- Homography can be represented by 3x3 non-singular matrix.
- Again, \mathbf{X} can project to the image plane. The projected point is $\mathbf{H}_\pi \mathbf{x}$, where \mathbf{H}_π is the homography from the image plane through plane π to another image plane.



Fundamental matrix



- $\mathbf{H}_\pi \mathbf{x}$ should be on the epipolar line \mathbf{l}' **whether $\mathbf{H}_\pi \mathbf{x}$ is not same with \mathbf{x}' .**

- Then, \mathbf{l}' can be written as


$$\mathbf{l}' = \mathbf{e}' \times \mathbf{H}_\pi \mathbf{x} = [\mathbf{e}']_\times \mathbf{H}_\pi \mathbf{x}$$

- The fundamental matrix \mathbf{F} is

$$\mathbf{F} = [\mathbf{e}']_\times \mathbf{H}_\pi$$

- The relation between \mathbf{x} and \mathbf{x}' is

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = \mathbf{x}'^T \mathbf{l}' = 0$$



The properties of fundamental matrix



- F is a rank 2 homogeneous matrix with 7 degrees of freedom.
- **Point correspondence:** If \mathbf{x} and \mathbf{x}' are corresponding image points, then $\mathbf{x}'^T F \mathbf{x} = 0$.
- **Epipolar lines:**
 - ◇ $\mathbf{l}' = F \mathbf{x}$ is the epipolar line corresponding to \mathbf{x} .
 - ◇ $\mathbf{l} = F^T \mathbf{x}'$ is the epipolar line corresponding to \mathbf{x}' .
- **Epipoles:**
 - ◇ $F \mathbf{e} = \mathbf{0}$.
 - ◇ $F^T \mathbf{e}' = \mathbf{0}$.
- **Computation from camera matrices P, P' :**
 - ◇ General cameras,
 $F = [\mathbf{e}']_{\times} P' P^+$, where P^+ is the pseudo-inverse of P , and $\mathbf{e}' = P' \mathbf{C}$, with $P \mathbf{C} = \mathbf{0}$.
 - ◇ Canonical cameras, $P = [I \mid \mathbf{0}]$, $P' = [M \mid \mathbf{m}]$,
 $F = [\mathbf{e}']_{\times} M = M^{-T} [\mathbf{e}]_{\times}$, where $\mathbf{e}' = \mathbf{m}$ and $\mathbf{e} = M^{-1} \mathbf{m}$.
 - ◇ Cameras not at infinity $P = K[I \mid \mathbf{0}]$, $P' = K'[R \mid \mathbf{t}]$,
 $F = K'^{-T} [\mathbf{t}]_{\times} R K^{-1} = [K' \mathbf{t}]_{\times} K' R K^{-1} = K'^{-T} R K^T [K R^T \mathbf{t}]_{\times}$.

Eight-point algorithm




We want to get a fundamental matrix from two images in different view.

In the images, at least 8 corresponding points are given.

Left image



Right image



Eight-point algorithm



We want to get a fundamental matrix from two images in different view.

In the images, at least 8 corresponding points are given.

If there are m correspondences, they satisfy

$$\mathbf{x}_i'^T \mathbf{F} \mathbf{x}_i = 0 \quad i = 1, \dots, m \quad \text{where} \quad \mathbf{x}_i = \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix} \quad \mathbf{x}_i' = \begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} \quad \mathbf{F} = \begin{bmatrix} f_{11} & f_{12} & f_{13} \\ f_{21} & f_{22} & f_{23} \\ f_{31} & f_{32} & f_{33} \end{bmatrix}$$

Eight-point algorithm



It can be represented by 9 unknown linear system.

$$\mathbf{A}\mathbf{f} = \mathbf{0}$$

where

$$\mathbf{A} = \begin{bmatrix} x'x & x'y & x' & y'x & y'y & y' & x & y & 1 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ x'x & x'y & x' & y'x & y'y & y' & x & y & 1 \end{bmatrix}$$

$$\mathbf{f} = \begin{bmatrix} f_{11} \\ f_{12} \\ f_{13} \\ f_{21} \\ f_{22} \\ f_{23} \\ f_{31} \\ f_{32} \\ f_{33} \end{bmatrix}$$

Eight-point algorithm



Details: $\mathbf{x}_i'^T \mathbf{F} \mathbf{x}_i = 0$ to $A\mathbf{f} = 0$

$$\mathbf{x}_i'^T \mathbf{F} \mathbf{x}_i = 0$$

Eight-point algorithm



Details: $\mathbf{x}_i'^T \mathbf{F} \mathbf{x}_i = 0$ to $\mathbf{A} \mathbf{f} = 0$

$$\begin{bmatrix} x'_i & y'_i & 1 \end{bmatrix} \begin{bmatrix} f_{11} & f_{12} & f_{13} \\ f_{21} & f_{22} & f_{23} \\ f_{31} & f_{32} & f_{33} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix} = 0$$

Eight-point algorithm



Details: $\mathbf{x}_i'^T \mathbf{F} \mathbf{x}_i = 0$ to $\mathbf{A} \mathbf{f} = 0$

$$\begin{aligned} & x_i' f_{11} x_i + x_i' f_{12} y_i + x_i' f_{13} 1 \\ & + y_i' f_{21} x_i + y_i' f_{22} y_i + y_i' f_{23} 1 \\ & + 1 f_{31} x_i + 1 f_{32} y_i + 1 f_{33} 1 = 0 \end{aligned}$$

Eight-point algorithm



Details: $\mathbf{x}_i'^T \mathbf{F} \mathbf{x}_i = 0$ to $\mathbf{A} \mathbf{f} = 0$

$$\begin{bmatrix} x_i'x_i & x_i'y_i & x_i' & y_i'x_i & y_i'y_i & y_i' & x_i & y_i & 1 \end{bmatrix} \begin{bmatrix} f_{11} \\ f_{12} \\ f_{13} \\ f_{21} \\ f_{22} \\ f_{23} \\ f_{31} \\ f_{32} \\ f_{33} \end{bmatrix} = 0$$

- Each equation for index i becomes each row of $\mathbf{A} \mathbf{f} = 0$.

Eight-point algorithm



The only nonzero solution of $\mathbf{A}\mathbf{x} = 0$ can exist if $\text{rank}(\mathbf{A}) = 9 - 1 = 8$
(“only” up to scaling \mathbf{f})

Each correspondence make one equation (a row of \mathbf{A})

It need eight points!

Eight-point algorithm: Implementation



1. Solve $\min_{\mathbf{f}} \|\mathbf{A}\mathbf{f}\|^2$ subject to $\|\mathbf{f}\|^2 = 1$
 - $\mathbf{f} \leftarrow$ the eigenvector of $\mathbf{A}^T \mathbf{A}$ corresponding the smallest eigenvalue.
2. $\mathbf{F}(3 \times 3 \text{ fundamental matrix}) \leftarrow \text{reshape}(\mathbf{f})$
3. Enforce \mathbf{F} to have rank 2
 - $\mathbf{F} = \mathbf{U}\mathbf{S}\mathbf{V}^T$ where \mathbf{U} , \mathbf{S} , \mathbf{V} are a singular value decomposition for \mathbf{F} .
Make the minimum singular value for \mathbf{S} become zero.
(the diagonal entries of \mathbf{S} are the singular values for \mathbf{S} .)
 - $\mathbf{F} \leftarrow \mathbf{U}\mathbf{S}\mathbf{V}^T$ by using modified \mathbf{S} at b.

Eight-point algorithm: Proof 1.



If there are more than 8 correspondences, we should get an approximation.

$$\min_{\mathbf{f}} \|\mathbf{A}\mathbf{f}\|^2 \quad \text{subject to} \quad \|\mathbf{f}\|^2 = 1$$

$$\mathbf{g}(\mathbf{f}) = \|\mathbf{A}\mathbf{f}\|^2 = (\mathbf{A}\mathbf{f})^T (\mathbf{A}\mathbf{f}) = \mathbf{f}^T \mathbf{A}^T \mathbf{A} \mathbf{f}$$

$$\mathbf{h}(\mathbf{f}) = 1 - \|\mathbf{f}\|^2 = 1 - \mathbf{f}^T \mathbf{f}$$

Eight-point algorithm: Proof 1.



Make the Lagrangian of the optimization.

$$L(\mathbf{f}, \lambda) = \mathbf{g}(\mathbf{f}) - \lambda \mathbf{h}(\mathbf{f}) = \mathbf{f}^T \mathbf{A}^T \mathbf{A} \mathbf{f} - \lambda (1 - \mathbf{f}^T \mathbf{f})$$

$$\min_{\mathbf{f}} \|\mathbf{A} \mathbf{f}\|^2$$

$$s.t. \quad \|\mathbf{f}\|^2 = 1$$



$$\min_{\mathbf{f}} L(\mathbf{f}, \lambda)$$

Eight-point algorithm: Proof 1.



Take derivatives of the Lagrangian.

$$\partial_{\mathbf{f}} L(\mathbf{f}, \lambda) = \mathbf{A}^T \mathbf{A} \mathbf{f} - \lambda \mathbf{f} = 0$$

$$\partial_{\lambda} L(\mathbf{f}, \lambda) = 1 - \mathbf{f}^T \mathbf{f} = 0$$

\mathbf{f} is normalized eigenvector of $\mathbf{A}^T \mathbf{A}$

Eight-point algorithm: Proof 1.



Let \mathbf{e}_λ is an eigenvector with eigenvalue λ .

$$\mathbf{g}(\mathbf{e}_\lambda) = \mathbf{e}_\lambda^T \mathbf{A}^T \mathbf{A} \mathbf{e}_\lambda = \mathbf{e}_\lambda^T \lambda \mathbf{e}_\lambda = \lambda$$

The eigenvector with the smallest eigenvalue is the result.

Eight-point algorithm: Proof 1.



$$\mathbf{g}(\mathbf{e}_\lambda) = \mathbf{e}_\lambda^T \mathbf{A}^T \mathbf{A} \mathbf{e}_\lambda = \mathbf{e}_\lambda^T \lambda \mathbf{e}_\lambda = \lambda$$

Details

$$\mathbf{A}^T \mathbf{A} \xrightarrow{\text{eigen decom.}} \mathbf{U} \mathbf{D} \mathbf{U}^{-1}$$

Numerically unstable

$$\begin{aligned} \mathbf{A} &\xrightarrow{\text{SVD}} \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T \\ &\Rightarrow \mathbf{A}^T \mathbf{A} = \mathbf{V} \mathbf{\Sigma}^2 \mathbf{V}^T \end{aligned}$$

✓ Preferable

- In general, eigen decomposition may produce complex values for real input matrix while singular value decomposition (SVD) always produces real values.
- The exact values of eigen decomposition of a real symmetric ($\mathbf{M}^T = \mathbf{M}$) matrix are real, but numerical methods produce little complex-valued error.
- For a symmetric matrix formed $\mathbf{A}^T \mathbf{A}$, utilizing SVD for \mathbf{A} is preferable.

Eight-point algorithm: Proof 3.



Is the result F have rank 2? \rightarrow It is not guaranteed.

We should reduce the dimension by singular value decomposition.

- Get SVD of F

$$\mathbf{F} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$$

$$\mathbf{U} = [\mathbf{u}_1 \quad \mathbf{u}_2 \quad \mathbf{u}_3] \quad \mathbf{V} = [\mathbf{v}_1 \quad \mathbf{v}_2 \quad \mathbf{v}_3] \quad \mathbf{\Sigma} = \begin{bmatrix} \sigma_1 & 0 & 0 \\ 0 & \sigma_2 & 0 \\ 0 & 0 & \sigma_3 \end{bmatrix}$$

- Set the smallest singular values to 0

$$\hat{\mathbf{\Sigma}} = \begin{bmatrix} \sigma_1 & 0 & 0 \\ 0 & \sigma_2 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

- Recompute F

$$\hat{\mathbf{F}} = \mathbf{U}\hat{\mathbf{\Sigma}}\mathbf{V}^T$$

Normalized eight-point algorithm



- Center the image data at the origin, and scale it so the mean squared distance between the origin and the data points is 2 pixels
“utils.py” , function “normalize_points”
- Use the eight-point algorithm to compute \mathbf{F} from the normalized points
- Enforce the rank-2 constraint (for example, take SVD of \mathbf{F} and throw out the smallest singular value)
- Transform fundamental matrix back to original units: if \mathbf{T} and \mathbf{T}' are the normalizing transformations in the two images, then the fundamental matrix in original coordinates is $\mathbf{T}'^T \mathbf{F} \mathbf{T}$

Rectification




- Left and right image should be reprojected onto the common plane parallel to the line between camera centers
➔ **Defined by homography matrices H and H'**
- Calculating H and H' when F and corresponding points are known is already implemented as the OpenCV native function.
- You need to slightly modify H and H' to avoid cropping and apply them to the left and right images.

Stereo image rectification

- Reproject image planes onto a common plane parallel to the line between camera centers
- Pixel motion is horizontal after this transformation
- Rectification: Two homographies (3x3 transforms), one for each input image reprojection

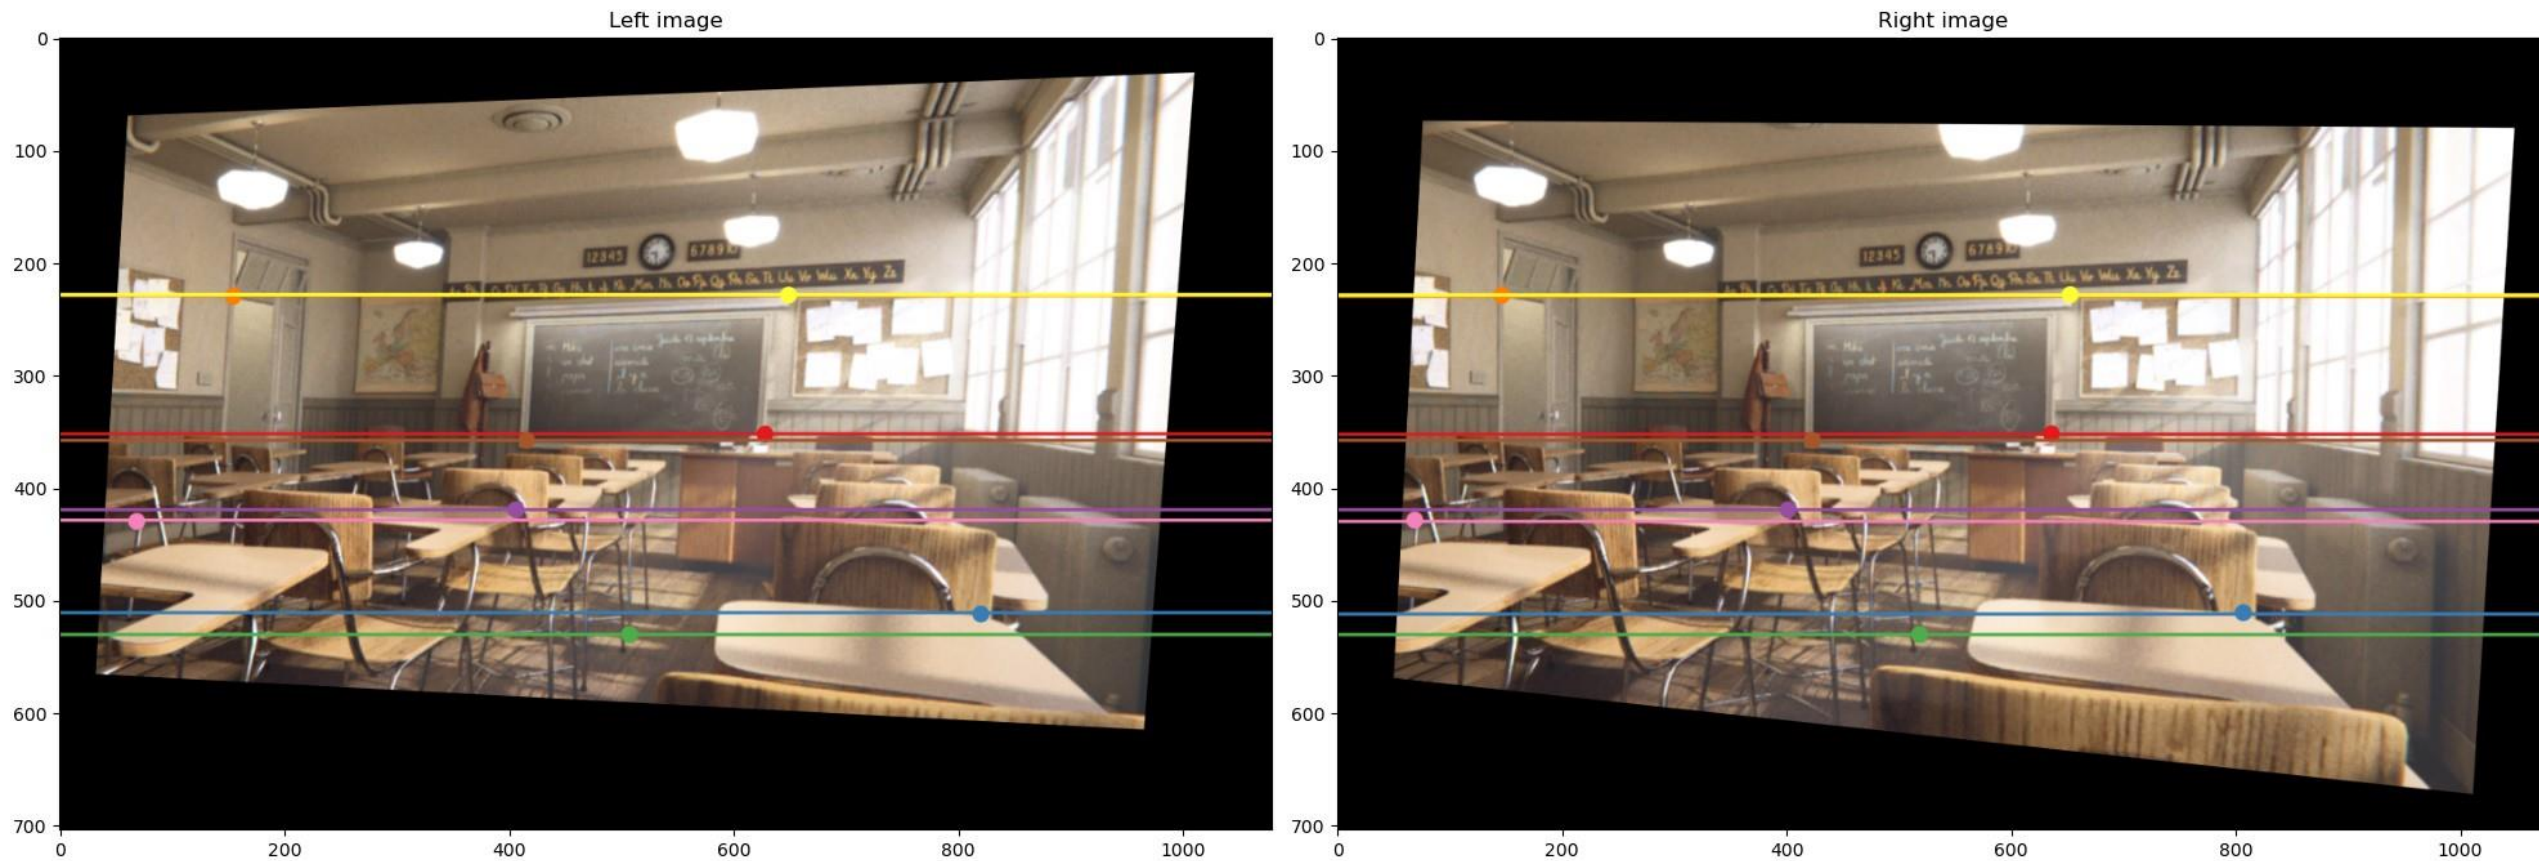
➤ C. Loop and Z. Zhang. [Computing Rectifying Homographies for Stereo Vision](#). IEEE Conf. Computer Vision and Pattern Recognition, 1999.

Lecture 1: Intro to Computer Vision CS484: Introduction to Computer Vision



Slide credit: Kristen Grauman

Rectification



result/rectified_imgs_epipolar_overlay.png
(generated in hw3_main.py after calling rectify_stereo_images)

Rectification



Before rectification



After rectification





- When you implement the part applying H and H' to the left and right images, following openCV functions might be helpful:
 - `cv2.warpPerspective()`
https://docs.opencv.org/4.5.0/da/d54/group__imgproc__transform.html#gaf73673a7e8e18ec6963e3774e6a94b87
 - `cv2.perspectiveTransform()`
https://docs.opencv.org/4.5.0/d2/de8/group__core__array.html#gad327659ac03e5fd6894b90025e6900a7
- Your rectified images should be aligned well as shown in the previous page.


Disparity map



- We provide perfect rectified image (img3.png (left), img4.png (right))
- And Ground truth disparity map of img3 (img3_disp.exr)



Disparity map



- This is one example of disparity map result.
 - It may poorly works in some region due to the inherent limitation of uncalibrated stereo problem.
 - Also it depends on your hyperparameter.
-
- Disparity map can be improved by using
 - Cost aggregation with box filter (+5 pts)
 - More sophisticated cost aggregation (+5 pts)
 - Sub-pixel disparity (+10 pts)
 - Calibrated cameras (not for this homework)

Evaluation of disparity map



Evaluation functions are already in “utils.py” and used in “hw3_main.py”

EPE (End-point-error)

$\text{epe} < 3.0$	15pts
$3.0 \leq \text{epe} < 4.0$	10pts
$4.0 \leq \text{epe} < 5.0$	5pts
$5.0 \leq \text{epe}$	0pts

Bad pixel ratio %


Ratio of pixels epe larger than 3.0

Bad pix $< 25\%$	15pts
$25\% \leq \text{Bad pix} < 30\%$	10pts
$30\% \leq \text{Bad pix} < 40\%$	5pts
$40\% \leq \text{Bad pix}$	0pts


(Optional) Sub-pixel disparity



Try this after you successfully implement disparity map!



- EPE: 2.4632
- Bad pixel ratio: 21.41%



- EPE: 2.3567
- Bad pixel ratio: 21.31%